

Distance Vector Routing Algorithm

iterative:

- continues until no nodes exchange info.
- *self-terminating*: no “signal” to stop

asynchronous:

- nodes need *not* exchange info/iterate in lock step!

distributed:


- each node communicates *only* with directly-attached neighbors

Distance Table data structure

- each node has its own
 - row for each possible destination
 - column for each directly-attached neighbor to node
- example: in node X, for dest. Y via neighbor Z:

$$\begin{aligned} D^X(Y,Z) &= \text{distance from X to Y, via Z as next hop} \\ &= c(X,Z) + \min_w \{D^Z(Y,w)\} \end{aligned}$$

Distance Table: example



$$D^E(C,D) = c(E,D) + \min_w \{D^D(C,w)\}$$

$$= 2+2 = 4$$

$$D^E(A,D) = c(E,D) + \min_w \{D^D(A,w)\}$$

$$= 2+3 = 5 \text{ loop!}$$

$$D^E(A,B) = c(E,B) + \min_w \{D^B(A,w)\}$$

$$= 8+6 = 14 \text{ loop!}$$

		cost to destination via		
destination	$D^E()$	A	B	D
	A			
	B			
	C			
	D			

Distance Table: example

Cost of going to **C** via **D**

$$\begin{aligned} D^E(C,D) &= c(E,D) + \min_w \{D^D(C,w)\} \\ &= 2+2 = 4 \end{aligned}$$

Cost of going to D via A

$$\begin{aligned} D^E(A,D) &= c(E,D) + \min_w \{D^D(A,w)\} \\ &= 2+3 = 5 \end{aligned}$$

loop!

$$\begin{aligned} D^E(A,B) &= c(E,B) + \min_w \{D^B(A,w)\} \\ &= 8+6 = 14 \end{aligned}$$


loop!

cost to destination via

$D^E()$	A	B	D
A	1	14	5
B	7	8	5
C	6	9	4
D	4	11	2

destination

Distance Table: example



		cost to destination via		
destination	$D^E()$	A	B	D
	A			
	B			
	C			
	D			

Distance table gives routing table

destination	cost to destination via			
	$D^E()$	A	B	D
	A	1	14	5
	B	7	8	5
	C	6	9	4
	D	4	11	2

destination	Outgoing link to use, cost	
	A	A,1
	B	D,5
	C	D,4
	D	D,2

Distance table \longrightarrow Routing table

Distance Vector Routing: overview

Each node:

Iterative, asynchronous: each

local iteration caused by:

- local link cost change
- message from neighbor: its least cost path change from neighbor

Distributed:


- each node notifies **neighbors only** when its least cost path to any destination changes
 - **neighbors then notify their neighbors if necessary**

wait for (change in local link cost or msg from neighbor)

recompute distance table

if least cost path to any dest has changed, *notify* neighbors

Distance Vector Algorithm: example



		cost via	
d	e s t	D ^X	
		Y	Z
Y	Z	2	∞
		∞	7

		cost via	
d	e s t	D ^Y	
		X	Z
X	Z	2	∞
		∞	1


		cost via	
d	e s t	D ^Z	
		X	Y
X	Y	7	∞
		∞	1

		cost via	
d	e s t	D ^X	
		Y	Z
Y	Z	2	8
		3	7

$$D^X(Z, Y) = c(X, Y) + \min_w \{D^Y(Z, w)\} \\ = 2 + 1 = 3$$

$$D^X(Y, Z) = c(X, Z) + \min_w \{D^Z(Y, w)\} \\ = 7 + 1 = 8$$

Distance Vector Algorithm: example



		cost via	
		Y	Z
dest	X		
	Y	2	∞
	Z	∞	7

		cost via	
		Y	Z
dest	X		
	Y	2	8
	Z	3	7

		cost via	
		Y	Z
dest	X		
	Y		
	Z		

		cost via	
		X	Z
dest	Y		
	X	2	∞
	Z	∞	1

		cost via	
		X	Z
dest	Y		
	X	2	8
	Z	9	1

		cost via	
		X	Z
dest	Y		
	X		
	Z		

		cost via	
		X	Y
dest	Z		
	X	7	∞
	Y	∞	1


		cost via	
		X	Y
dest	Z		
	X	7	3
	Y	9	1

		cost via	
		X	Y
dest	Z		
	X		
	Y		


Link cost changes

Link cost changes:

- node detects local link cost change
- updates distance table
- if cost change in least cost path, notify neighbors



“good
news
travels
fast”





algorithm
terminates

Link cost changes

Link cost changes:


- good news travels fast
- bad news travels slow - “count to infinity” problem!



Link cost changes

Link cost changes:

- good news travels fast
- bad news travels slow - “count to infinity” problem!





algorithm
continues
on!

- When will it STOP?

DV: poisoned reverse

If Z routes through Y to get to X :

- Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)
- will this completely solve count to infinity problem?



Comparison of LS and DV algorithms

Message complexity

- LS: Each node needs to know the cost of each link (n nodes, E links) $O(nE)$ messages
- DV: exchange between neighbors only
 - convergence time varies

Speed of Convergence

- LS: $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- DV: convergence time varies
 - may be routing loops
 - count-to-infinity problem

Robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its *own* table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate thru network

Making routing scalable

our routing study thus far - idealized

- all routers identical
- network “flat”

... *not* true in practice

scale: with billions of destinations:

- can't store all destinations in routing tables!
- routing table exchange would swamp links!

administrative autonomy

- internet = network of networks
- each network admin may want to control routing in its own network

Internet approach to scalable routing

aggregate routers into regions known as “**autonomous systems**” (AS) (a.k.a. “domains”)

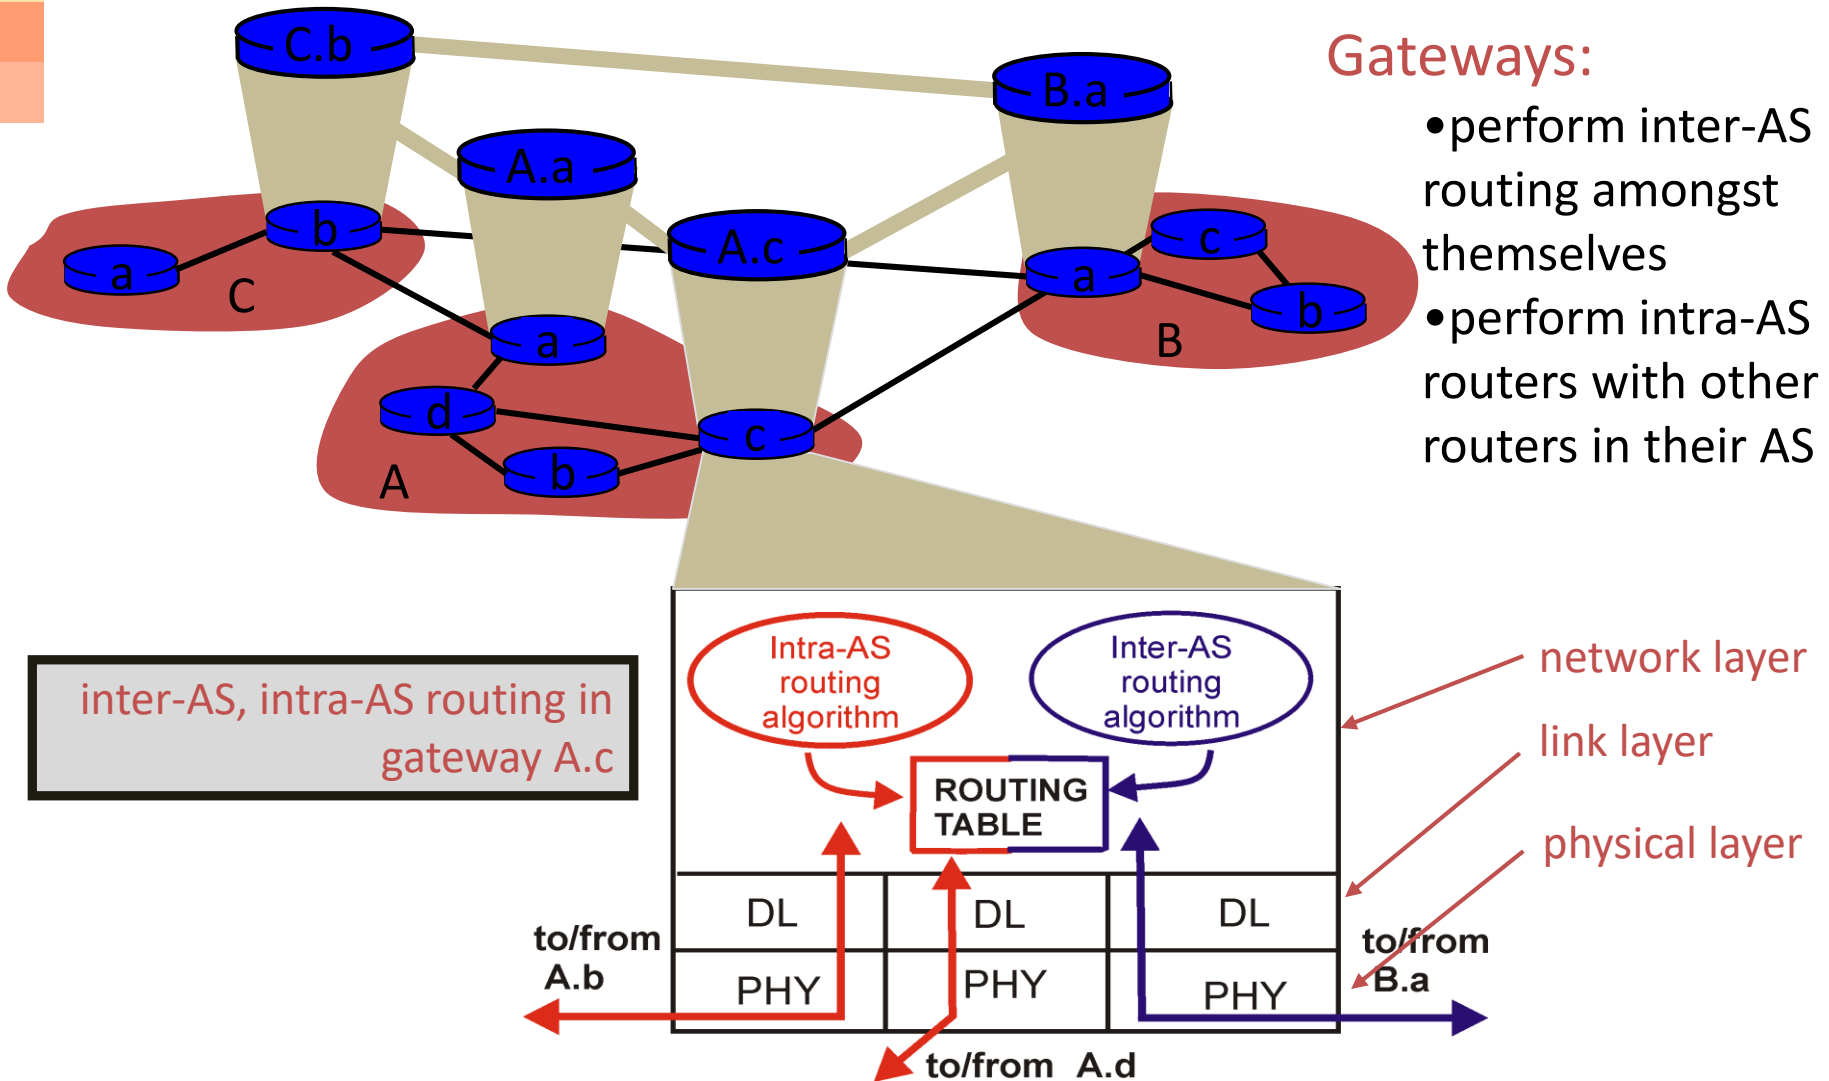
intra-AS routing

- routing among hosts, routers in same AS (“network”)
- all routers in AS must run *same* intra-domain protocol
- routers in *different* AS can run *different* intra-domain routing protocol
- gateway router: at “edge” of its own AS, has link(s) to router(s) in other AS'es

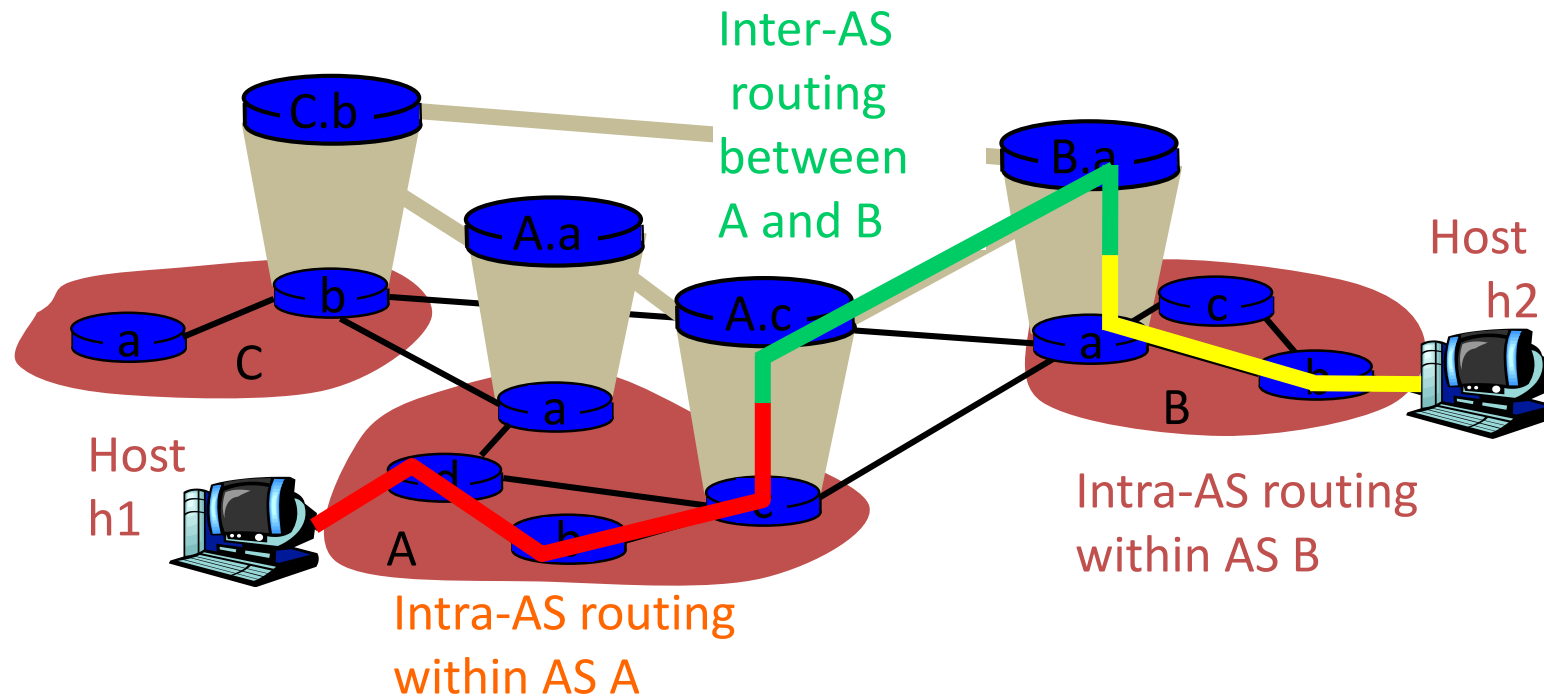
inter-AS routing

- routing among AS'es
- gateways perform inter-domain routing (as well as intra-domain routing)

Intra-AS and Inter-AS routing



Intra-AS and Inter-AS routing



- Intra-AS protocols include RIP (DV), EIGRP (Hybrid) and OSPF (LSA).
- Inter-AS protocols include BGP (Path Vector Protocol). BGP lets every AS in the Internet know about the subnets in YOUR AS.